

AI DARK FACTORY (HACKATHON EDITION)

tablekeeper generator

Track: 🍴 tablekeeper



THE BAND

Four agents, one room



Planner

Plans each stage



Builder

Writes the code



Tester

Tests against the spec



Delivery

Ships a clean stage

WHO DOES WHAT



Planner

The lead. Turns each stage's specification into a step-by-step plan, writes fix plans from bug reports, confirms deliveries and starts the next stage. Never writes code.

TALKS TO

Builder, Delivery



Builder

Builds the app from the plan, runs it in the sandbox, checks it against the plan and commits. Never decides its own work is done.

TALKS TO

Tester



Tester

Spawns two or three short-lived tester agents, one per kind of user, and runs the official checker. Reads the code but can never change it.

TALKS TO

**Planner (bugs)
Delivery (pass)**



Delivery

Turns a passing stage into a clean, documented folder, scans it for secrets and confirms it still builds. Never changes behavior.

TALKS TO

Planner

HOW A STAGE MOVES

One loop per stage



Bugs found: Tester reports to Planner, who writes a fix plan for Builder

Delivered: Planner sends Delivery a receipt and starts the next stage

Builder never talks to Delivery

Only human input: one prompt

SETUP

How the factory runs

AGENT FRAMEWORK

Hermes Agent v2026.9.24

with BAND's Hermes plugin, one agent per seat

MODEL

deepseek/deepseek-v4.1-flash

the same model for all four agents

PROVIDER

OpenRouter

Featherless works as an alternative

SANDBOX

One Docker container per agent

no internet except an allowlist, private build daemon

CATCHING BAD WORK

Stage 1: rejected, fixed, passed



Tester

Round 1: **FAIL**. Five bugs: the data reset accepted bad ids and references, and repeat-request protection ignored the request path.



Planner

Reads the report next to the builder's code and writes one fix plan for all five.



Builder

Fixes all five and hands the new version back to the tester.



Tester

Round 2: **PASS**, with the earlier checks still green.

The tester also rejected stages 3 and 4 once each

Every report, fix plan and fix is its own commit

RESULT

All four stages, one prompt

STAGES BUILT

4 of 4

each passes the official
checker

TIME

2 h 5 min

prompt to final report

MODEL COST

\$2.82

measured on OpenRouter

BUGS CAUGHT

3 stages

rejected, fixed, then
passed

A restaurant booking app like OpenTable: search and an availability grid, sign up and log in, bookings that are never double-booked, combined tables, booking policies, history, recurring bookings and manager replans.